

Generative Modeling by Estimating Gradients of the Data Distribution

Based on

<https://yang-song.net/blog/2021/score/>

Song et al., 2019 (NeurIPS)

Overview

- Blog-post inspired (quite analogous to the paper)
 - Includes some ‘generative modeling’ recap
- I’ll go fast through the overall idea to get a full picture
 - We can the clarify details
- Triple-check me



Objective of generative modeling

dataset $\{\mathbf{x}_1, \mathbf{x}_2, \dots, \mathbf{x}_N\}$, underlying data distribution $p(\mathbf{x})$.

Objective: Learn $p_\theta(\mathbf{x})$ (with option to sample)

Energy-based:

$$p_\theta(\mathbf{x}) = \frac{e^{-f_\theta(\mathbf{x})}}{Z_\theta},$$

Objective: maximum log likelihood

$$\max_{\theta} \sum_{i=1}^N \log p_\theta(\mathbf{x}_i).$$

$f_\theta(\mathbf{x})$ is often called an unnormalized probabilistic model, or energy-based model [7].

Neural net predicts energy of a given state (sample).

Good: We have a formulation to train a model of our data distribution based on maximum log likelihood

Complications:

1. Z_θ intractable
2. How to sample from f_θ ?

One solution

$$p_{\theta}(\mathbf{x}) = \frac{e^{-f_{\theta}(\mathbf{x})}}{Z_{\theta}},$$

intractable



- Invertible models (last week)
 - $z = f_{\theta}(x)$
 - $p(z) = \mathcal{N}(0, 1; z)$
 - $p_{\theta}(x) = p(z) * |\det df(x)/dx|$

New way (today)

$$p_{\theta}(\mathbf{x}) = \frac{e^{-f_{\theta}(\mathbf{x})}}{Z_{\theta}}, \quad \text{intractable}$$

$\nabla_{\mathbf{x}} \log p(\mathbf{x})$ is called: (Stein) score function

$$\nabla_{\mathbf{x}} \log p_{\theta}(\mathbf{x}) = -\nabla_{\mathbf{x}} f_{\theta}(\mathbf{x}) - \underbrace{\nabla_{\mathbf{x}} \log Z_{\theta}}_{=0} = -\nabla_{\mathbf{x}} f_{\theta}(\mathbf{x})$$



Remaining challenges:


- How to train this (no more log likelihood maximization)
- How to sample from it

How to learn the score function?

Minimize Fisher divergence:

$$\mathbb{E}_{p(\mathbf{x})} [\|\nabla_{\mathbf{x}} \log p(\mathbf{x}) - \mathbf{s}_{\theta}(\mathbf{x})\|_2^2]$$

Unknown ☹️



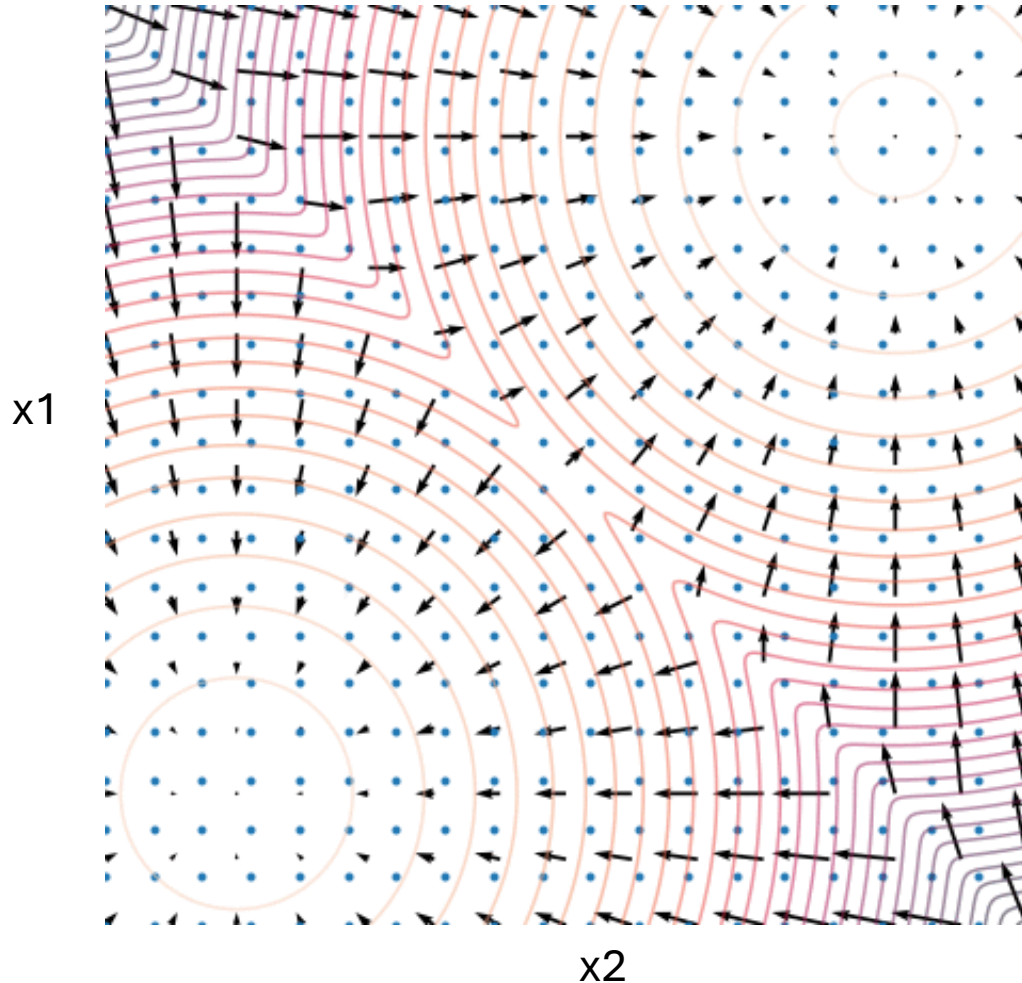
there exists a family of methods called **score matching**³ [16, 17, 31] that minimize the Fisher divergence without knowledge of the ground-truth data score. Score matching objectives can

“Score matching objectives”: Objective that leads to a match between scores of two distributions ($p(x)$, $p_{\theta}(x)$)

Postponed: Appendix 1

How to sample?

We just trained this




$$\mathbf{x}_0 \sim \pi(\mathbf{x}),$$

$$\mathbf{x}_{i+1} \leftarrow \mathbf{x}_i + \epsilon \nabla_{\mathbf{x}_i} \log p(\mathbf{x}) + \sqrt{2\epsilon} \mathbf{z}_i, \quad i = 0, 1, \dots, K,$$


Provably converges to a sample from $p(\mathbf{x})$ (if $\epsilon \rightarrow 0$ and $K \rightarrow \infty$)

Langevin (Monte Carlo) sampling

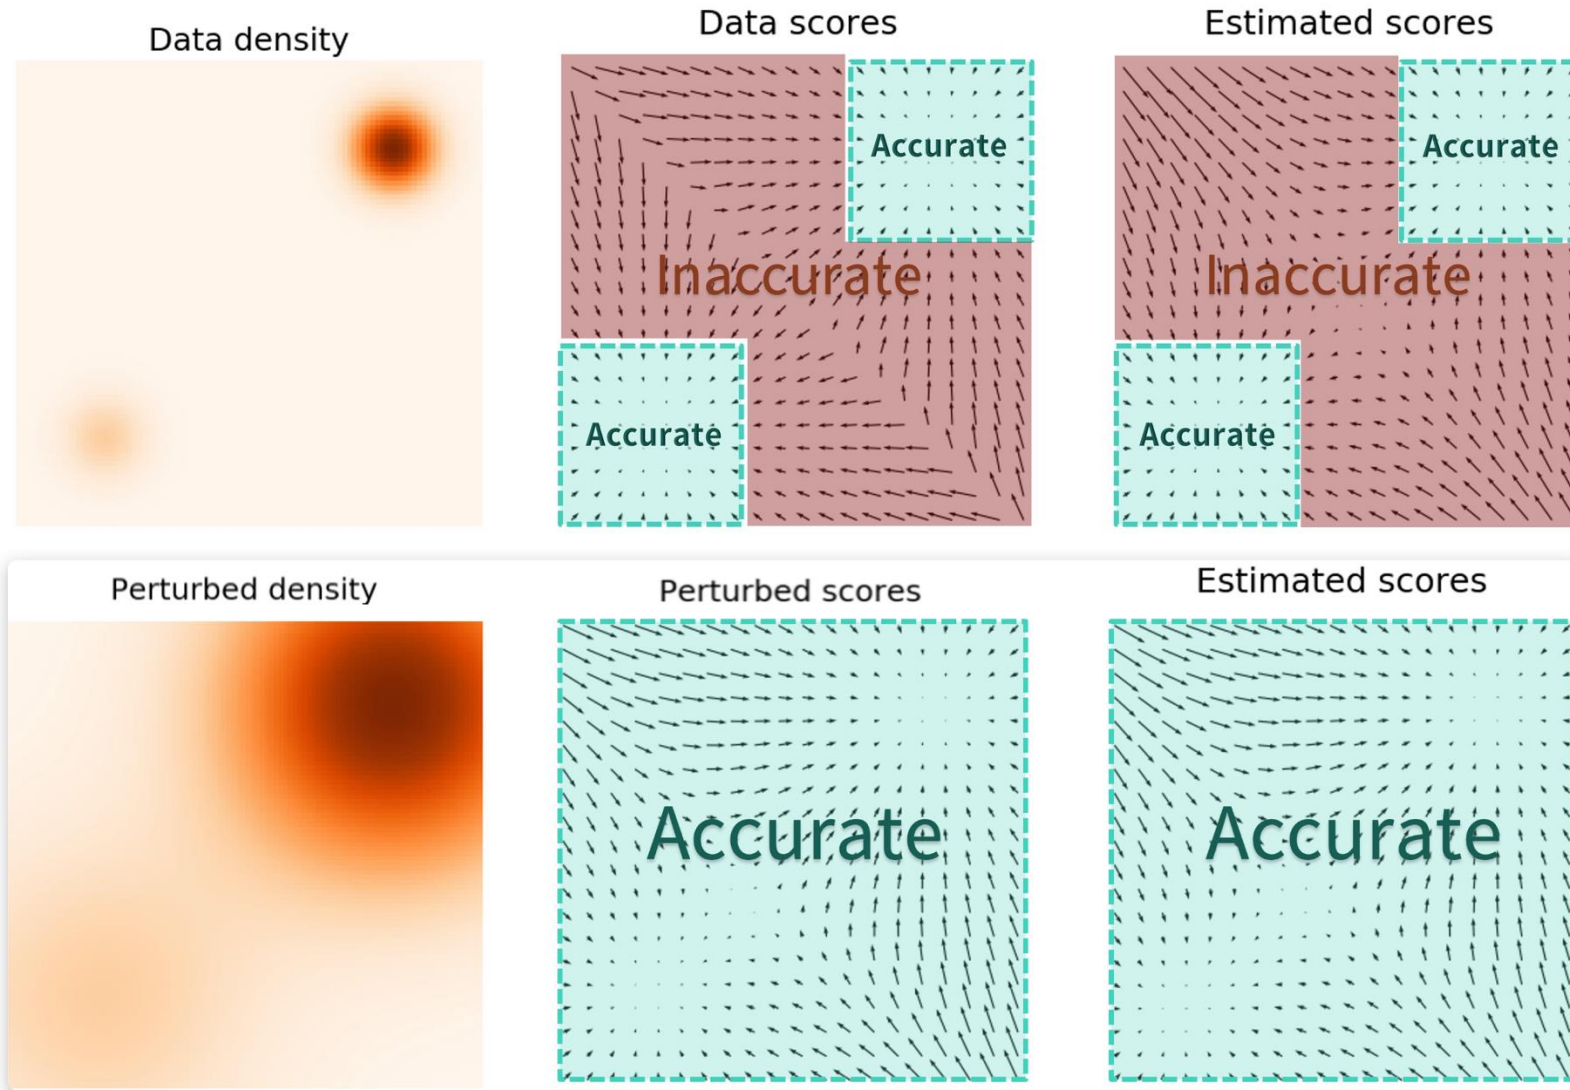
Challenge: Low density regions -> poor score estimates



$$\mathbb{E}_{p(\mathbf{x})} [\|\nabla_{\mathbf{x}} \log p(\mathbf{x}) - \mathbf{s}_\theta(\mathbf{x})\|_2^2] = \int p(\mathbf{x}) \|\nabla_{\mathbf{x}} \log p(\mathbf{x}) - \mathbf{s}_\theta(\mathbf{x})\|_2^2 d\mathbf{x}.$$



Noise to the rescue: Intuition




Simply add noise to samples and model

How to add noise?

- Too much: no signal
- Too little: poor space coverage

Note the 'i' in $s_{\theta}(\mathbf{x}, i)$

Noise Conditional Score-Based Model



$$p_{\sigma_i}(\mathbf{x}) = \int p(\mathbf{y}) \mathcal{N}(\mathbf{x}; \mathbf{y}, \sigma_i^2 I) d\mathbf{y}.$$


Note that we can easily draw samples from $p_{\sigma_i}(\mathbf{x})$ by sampling $\mathbf{x} \sim p(\mathbf{x})$ and computing $\mathbf{x} + \sigma_i \mathbf{z}$, with $\mathbf{z} \sim \mathcal{N}(0, I)$.



train $\mathbf{s}_{\theta}(\mathbf{x}, i) \approx \nabla_{\mathbf{x}} \log p_{\sigma_i}(\mathbf{x})$ for all $i = 1, 2, \dots, L$.

min $\sum_{i=1}^L \lambda(i) \mathbb{E}_{p_{\sigma_i}(\mathbf{x})} [\|\nabla_{\mathbf{x}} \log p_{\sigma_i}(\mathbf{x}) - \mathbf{s}_{\theta}(\mathbf{x}, i)\|_2^2],$ (sum of weighted Fishers) => compatible with score matching


Sampling from Noise Conditional Score-Based Model



- σ_3 model is good at the beginning (but poor final result)


$$\mathbf{x}_0 \sim \pi(\mathbf{x}),$$

- σ_1 model is good to refine



Recap

- Model score instead of $p(\mathbf{x})$
- Training: "score matching objective" $\mathbb{E}_{p(\mathbf{x})} [\|\nabla_{\mathbf{x}} \log p(\mathbf{x}) - \mathbf{s}_{\theta}(\mathbf{x})\|_2^2]$
- Sampling: Langevin $\mathbf{x}_{i+1} \leftarrow \mathbf{x}_i + \epsilon \nabla_{\mathbf{x}} \log p(\mathbf{x}) + \sqrt{2\epsilon} \mathbf{z}_i, \quad i = 0, 1, \dots, K,$
- Circumvent low density: Add noise



Appendix: Score matching objectives

Minimize Fisher divergence:

$$\mathbb{E}_{p(\mathbf{x})} [\|\nabla_{\mathbf{x}} \log p(\mathbf{x}) - \mathbf{s}_{\theta}(\mathbf{x})\|_2^2]$$

Unknown :(☹

Equivalent (up to a constant):

$$\mathbb{E}_{p_{\text{data}}(\mathbf{x})} \left[\text{tr}(\nabla_{\mathbf{x}} \mathbf{s}_{\theta}(\mathbf{x})) + \frac{1}{2} \|\mathbf{s}_{\theta}(\mathbf{x})\|_2^2 \right]$$

Equivalent to "Denoising score matching":

$$\frac{1}{2} \mathbb{E}_{\underbrace{q_{\sigma}(\tilde{\mathbf{x}}|\mathbf{x})}_{\text{Denoising}} \underbrace{p_{\text{data}}(\mathbf{x})}} [\|\mathbf{s}_{\theta}(\tilde{\mathbf{x}}) - \nabla_{\tilde{\mathbf{x}}} \log q_{\sigma}(\tilde{\mathbf{x}} | \mathbf{x})\|_2^2]$$


Connection to other types of models

Diffusion models

Same model family. Unifying framework: Song 2021 ICLR, Ho 2020

	Score-based model	Diffusion probabilistic model
Perturbation	Multiple scales of noise	Multiple scales of noise
Training objective	Score matching	ELBO
Sampling	Langevin (MCMC)	Learned decoder
Unique ability	Calculate log-likelihoods exactly	Can be made faster?

Generative stochastic networks



MCMC sampling, Denoising model