

Residual Flows for Invertible Generative Modeling

Ricky T. Q. Chen, Jens Behrmann, David Duvenaud,
Jörn-Henrik Jacobsen

Motivation

- Flow based generative models aim to invert a transformation (*noising*) and trained with ML
- iResnet offers invertibility directly through flexible architecture
- Calculating exact log likelihood requires trace of Jacobian, proposed estimation results in biased estimator
- They proposed a memory and computation efficient unbiased estimator for Jacobian log-det
 - Do not require ODE solvers unlike FM
 - Can directly calculate the likelihood -> Stochastic estimate of Jacobian log-det
 - iResnets are discrete and inverse requires iteration (*simulation in CNF*)
 - Unlike CNF, calculation of network's Lipschitz constant

Recap

- iResNet

- Forward pass $y = f(x) = x + g(x)$ remains bijective, x is recoverable from y
 - g is the NN part of the residual block
- Requires Lipschitz continuity of $<:1$
 - There is a guaranteed unique inverse \rightarrow contractiveness ensures inversion leads to a fixed point
- Process can be reverted through fixed point iteration (do the inversion n times)

$$x^{(i)} = y - g(x^{(i-1)})$$

- Reverting requires calculation the trace of Jacobian
 - Also, ensuring Lipschitz constraint on g requires calculation of it as well

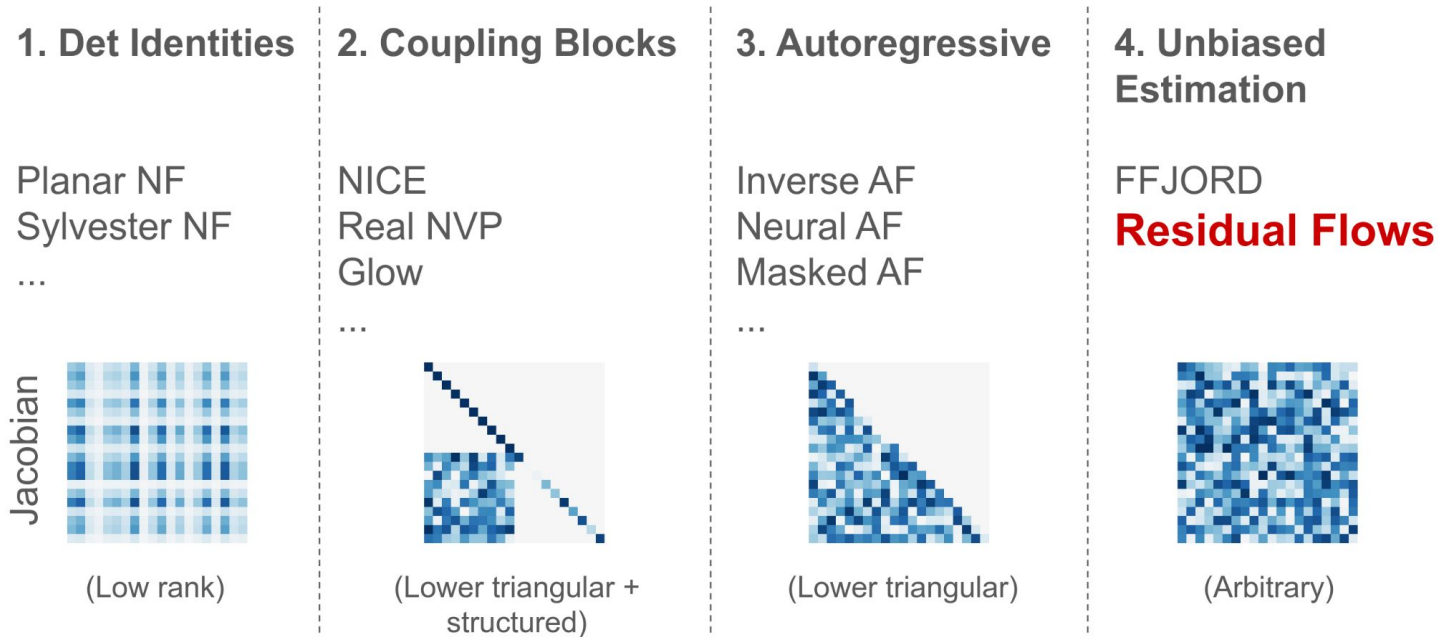
$$\|J_g(x)\|_2 = \|W_L \dots W_2 \phi'(z_1) W_1 \phi'(z_2)\|_2$$

- Can be used for estimating a tractable density

$$\mathcal{L}_{\text{NLL}} = -\frac{1}{N} \sum_{i=1}^N \left[\log p_Z(f(x_i)) + \log \left| \det \frac{\partial f}{\partial x_i} \right| \right] \text{ where } p_Z \text{ is a gaussian}$$

Problem of existing Flow Approaches

- They rely on specific architectures when predicting Jacobian during density estimation
- Results in **biased estimator, slow inference**



- Model can overfit on the bias without maxing ML
- Less flexibility in the architecture

Change of Variables for Estimating Density

- For an invertible transformation F , the density estimation:

$$\ln p_x(x) = \ln p_z(z) + \ln |\det J_F(x)|,$$

- Given that Residual Network $f(x)$:

$$y = f(x) = x + g(x).$$

- Det of Jacobian can be written as a power series from a book I don't have access to

$$\log p(x) = \log p(f(x)) + \text{tr} \left(\sum_{k=1}^{\infty} \frac{(-1)^{k+1}}{k} [J_g(x)]^k \right)$$

- The trace can be estimated by Skilling-Hutchinson estimator

$\text{tr}(A) = \mathbb{E}_{p(v)} [v^T A v]$, where v is a random variable (generally normal dist)

$$\mathbb{E}_v \left[\sum_{k=1}^{\infty} \frac{(-1)^{k+1}}{k} v^T [J_g(x)]^k v \right]$$

- The SH estimator is Lipschitz bounded as well (See App. proof. Theorem 1)

- Change of Variables Theorem

- In Bishop:

$$p_y(y) = p_x(x) \left| \frac{dx}{dy} \right|$$

- More Formally

$$\int_{\varphi(U)} f(v) dv = \int_U f(\varphi(u)) |\det(D\varphi)(u)| du.$$

- Change of Variables in flows

$$\log p(x) = \log p(f(x)) + \log \left| \det \frac{df(x)}{dx} \right|$$

<Model distribution> <Base distribution>

How to avoid calculating infinite series

- Truncation:
 - Just calculate first n terms

$$\mathbb{E}_v \left[\sum_{k=1}^{\infty} \frac{(-1)^{k+1}}{k} v^T [J_g(x)]^k v \right] \approx \mathbb{E}_v \left[\sum_{k=1}^n \frac{(-1)^{k+1}}{k} v^T [J_g(x)]^k v \right]$$

for Each residual block **do**

Lip constraint: $\hat{W}_j := \text{SN}(W_j, x)$ for linear Layer W_j .

Draw v from $\mathcal{N}(0, I)$

$w^T := v^T$

$\ln \det := 0$

for $k = 1$ **to** n **do**

$w^T := w^T J_g$ (vector-Jacobian product)

$\ln \det := \ln \det + (-1)^{k+1} w^T v / k$

end for

return

- Results in biased estimator which requires tradeoff

- Reduce d to reduce the Lip const
- Increase n to diminish the bias term
- Also model can overfit on on bias

$$\underbrace{\mathbb{E}_v \left[\sum_{k=1}^n \frac{(-1)^{k+1}}{k} v^T [J_g(x)]^k v \right]}_{\text{biased estimator}} + \underbrace{\text{tr} \left(\sum_{k=n+1}^{\infty} \frac{(-1)^{k+1}}{k} [J_g(x)]^k \right)}_{\text{bias}} \in \mathcal{O} \left(\frac{d}{1 - \text{Lip}(g)} \times \text{Lip}(g)^n \right)$$

How to avoid calculating infinite series

- With an **unbiased estimator**:

- Calculate the first term of the series Δ_1
- Flip a coin for the next terms, and weight the the calculated terms:

$$\mathbb{E} \left[\Delta_1 + \left[\frac{1}{1-q} \sum_{k=2}^{\infty} \Delta_k \right] \mathbb{1}_{b=0} + [0] \mathbb{1}_{b=1} \right]$$

- Has a probability of running finite time with probability of q
- It is an unbiased estimator because:

$$= \Delta_1 + \left[\frac{1}{1-q} \sum_{k=2}^{\infty} \Delta_k \right] (1 - q)$$

$$= \sum_{k=1}^{\infty} \Delta_k$$

How to avoid calculating infinite series


- With an **unbiased estimator**:
 - Modify it so that It can run in finite terms with $q=1$
 - Apply the same trick for each next term (*still an unbiased estimator*)

$$\sum_{k=1}^{\infty} \Delta_k = \mathbb{E}_{n \sim p(N)} \left[\sum_{k=1}^n \frac{\Delta_k}{\mathbb{P}(N \geq k)} \right]$$

kth term is weighted by
prob of seeing $\geq k$ tosses

- Final LLH: $\log p(x) = \log p(f(x)) + \mathbb{E}_{n,v} \left[\sum_{k=1}^n \frac{(-1)^{k+1}}{k} \frac{v^T [J_g(x)]^k v}{\mathbb{P}(N \geq k)} \right]$
 - Two levels of stochasticity:
 - Coin toss
 - Skilling-Hutchinson estimator

Calculate first 2 terms, then
sample from $\text{Geom}(0.5)$.



Engineering Problems

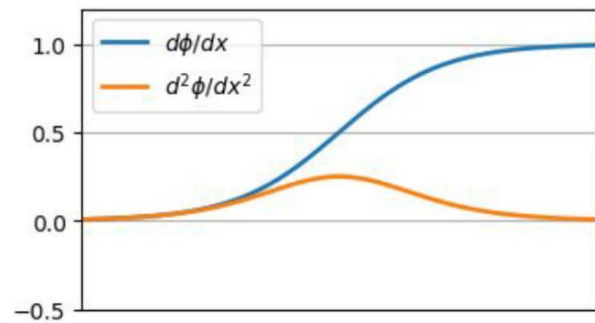
- OOM error further into the training:
 - # calculated terms could get high during training
 - Terms need to be backpropagated $\mathbb{E}_{n,v} \left[\sum_{k=1}^n \alpha_k \frac{\partial v^T [J_g(x)]^k v}{\partial \theta} \right]$
 - Use Neumann gradient series (?) to take the differentiation out of the power-series
 - Number of derivatives stored in memory becomes independent of n

$$\mathbb{E}_{n,v} \left[\left(\sum_{k=1}^n \alpha_k v^T [J_g(x)]^k \right) \frac{\partial J_g(x) v}{\partial \theta} \right] \longrightarrow \text{Only work if the series is finite}$$


Engineering Problems

- Lipschitz constraint requires regularization, such as spectral norm
- Lipschitz constrained activation functions can have a saturation in their derivatives
 - Manifest as second derivatives vanish
 - Swish is a good candidate, but breaks the uniform Lipschitz constraint
 - Scale it!


$$\text{LipSwish}(x) = \text{Swish}(x)/1.1$$



Softplus



ELU



LipSwish

Thank you for listening!